


# Gaussian Mixture Model (GMM) and Hidden Markov Model (HMM)

Samudravijaya K

Tata Institute of Fundamental Research, Mumbai

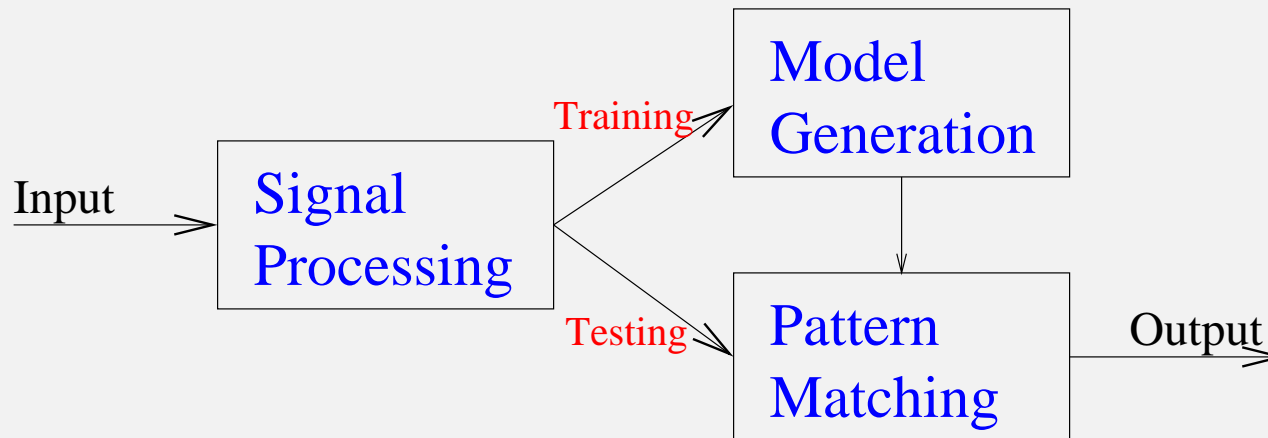
[chief@tifr.res.in](mailto:chief@tifr.res.in)



09-JAN-2009

Majority of the slides are taken from S.Umesh's tutorial on ASR (WiSSAP 2006).

# Pattern Recognition



GMM: static patterns

HMM: sequential patterns

# Basic Probability

## Joint and Conditional probability

$$p(A, B) = p(A|B) p(B) = p(B|A) p(A)$$

## Bayes' rule

$$p(A|B) = \frac{p(B|A) p(A)}{p(B)}$$

If  $A_i$ s are mutually exclusive events,

$$p(B) = \sum_i p(B|A_i) p(A_i)$$

$$p(A|B) = \frac{p(B|A) p(A)}{\sum_i p(B|A_i) p(A_i)}$$

# Normal Distribution

Many phenomenon are described by Gaussian *pdf*

$$p(x|\boldsymbol{\theta}) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{1}{2\sigma^2}(x - \mu)^2\right) \quad (1)$$


*pdf* is parameterised by  $\boldsymbol{\theta} = [\mu, \sigma^2]$  where mean =  $\mu$  and variance= $\sigma^2$ .

A convenient *pdf*: second order statistics is sufficient.

**Example:** Heights of Pygmies  $\Rightarrow$  Gaussian *pdf* with  $\mu = 4ft$  & std-dev( $\sigma$ ) =  $1ft$

**OR:** Heights of bushmen  $\Rightarrow$  Gaussian *pdf* with  $\mu = 6ft$  & std-dev( $\sigma$ ) =  $1ft$

**Question:** If we arbitrarily pick a person from a population  $\Rightarrow$   
what is the probability of the height being a particular value?




If I pick arbitrarily a Pygmy, say  $x$ , then

$$\Pr(\text{Height of } x=4'1'') = \frac{1}{\sqrt{2\pi \cdot 1}} \exp \left( -\frac{1}{2 \cdot 1} (4'1'' - 4)^2 \right) \quad (2)$$

*Note:* Here mean and variances are fixed, only the observations,  $x$ , change.

Also see:  $\Pr(x = 4'1'') \gg \Pr(x = 5')$     **AND**     $\Pr(x = 4'1'') \gg \Pr(x = 3')$



Conversely: Given a person's height is 4'1"  $\Rightarrow$

Person is **more likely** to be a pygmy than bushman.

If we observe heights of many persons – say 3'6", 4'1", 3'8", 4'5", 4'7", 4', 6'5" *and* all are from *same* population (i.e. either pygmy or bushmen.)

$\Rightarrow$  then more certain we are that the population is pygmy.

More the observations  $\Rightarrow$  better will be our decision

# Likelihood Function

$x[0], x[1], \dots, x[N-1]$

$\Rightarrow$  set of independent observations from *pdf* parameterised by  $\theta$ .

*Previous Example:*  $x[0], x[1], \dots, x[N-1]$  are heights observed and  $\theta$  is the mean of density which is unknown ( $\sigma^2$  assumed known).

$$\begin{aligned} L(\mathbf{X}; \theta) = p(x_0 \dots x_{N-1}; \theta) &= \prod_{i=0}^N p(x_i; \theta) \\ &= \frac{1}{(2\pi\sigma^2)^{N/2}} \exp \left( -\frac{1}{2\sigma^2} \sum_{i=0}^N (x_i - \theta)^2 \right) \quad (3) \end{aligned}$$


$L(\mathbf{X}; \theta)$  is a function of  $\theta$  and is called **Likelihood Function**

**Given:**  $x_0 \dots x_{N-1}$ ,  $\Rightarrow$  what can we say about value of  $\theta$ , i.e. best estimate of  $\theta$ .

# Maximum Likelihood Estimation

**Example:** We know height of a person  $x[0] = 4'4''$ .

Most likely to have come from which *pdf*  $\Rightarrow \theta = 3', 4'6''$  or  $6'$  ?



Maximum of  $L(x[0]; \theta = 3')$ ,  $L(x[0]; 4'6'')$  and  $L(x[0]; \theta = 6')$   $\Rightarrow$  choose  $\hat{\theta} = 4'6''$ .

If  $\theta$  is just a parameter, we will choose  $\arg \max_{\theta} L(x[0]; \theta)$ .



# Maximum Likelihood Estimator

Given  $x[0], x[1], \dots, x[N-1]$  and *pdf* parameterised by  $\boldsymbol{\theta} = \begin{bmatrix} \theta_1 \\ \theta_2 \\ \cdot \\ \cdot \\ \theta_{m-1} \end{bmatrix}$

We form Likelihood function  $L(\mathbf{X}; \boldsymbol{\theta}) = \prod_{i=0}^N p(x_i; \boldsymbol{\theta})$

$$\hat{\boldsymbol{\theta}}_{MLE} = \arg \max_{\boldsymbol{\theta}} L(\mathbf{X}; \boldsymbol{\theta})$$

For height problem:

$\Rightarrow$  can show  $(\hat{\theta})_{MLE} = \frac{1}{N} \sum x_i$

$\Rightarrow$  Estimate of mean of Gaussian = sample mean of measured heights.

# Bayesian Estimation

- **MLE**  $\Rightarrow \theta$  is assumed unknown but deterministic
- **Bayesian Approach**:  $\theta$  is assumed random with pdf  $p(\theta) \Rightarrow$  Prior Knowledge.

$$\underbrace{p(\theta|\mathbf{x})}_{\text{Aposterior}} = \frac{p(\mathbf{x}|\theta)p(\theta)}{p(\mathbf{x})} \propto p(\mathbf{x}|\theta) \underbrace{p(\theta)}_{\text{Prior}}$$


- **Height problem**: Unknown mean is random  $\Rightarrow$  pdf Gaussian  $\mathcal{N}(\gamma, \nu^2)$

$$p(\mu) = \frac{1}{\sqrt{2\pi\nu^2}} \exp\left(-\frac{1}{2\nu^2}(\mu - \gamma)^2\right)$$

$$\text{Then : } (\hat{\mu})_{\text{Bayesian}} = \frac{\sigma^2\gamma + n\nu^2\bar{x}}{\sigma^2 + n\nu^2}$$

$\Rightarrow$  Weighted average of sample mean and a *prior* mean

# Gaussian Mixture Model




$$p(x) = \alpha p(x|N(\mu_1; \sigma_1)) + (1 - \alpha) p(x|N(\mu_2; \sigma_2))$$

$$p(x) = \sum_{m=1}^M w_m p(x|N(\mu_m; \sigma_m)), \quad \sum w_i = 1$$

## Characteristics of GMM:

Just like ANNs are universal approximators of functions, GMMs are universal approximators of densities (provided sufficient no. of mixtures are used); true for diagonal GMMs as well.


# General Assumption in GMM



- Assume that there are  $M$  components.
- Each component generates data from a Gaussian with mean  $\mu_m$  and covariance matrix  $\Sigma_m$ .


# GMM

Consider the following probability density function shown in solid blue



It is useful to parameterise or “model” this seemingly arbitrary “blue” *pdf*

## Gaussian Mixture Model (Contd.)



Actually – *pdf* is a mixture of 3 Gaussians, i.e.

$$p(x) = c_1 N(x; \mu_1, \sigma_1) + c_2 N(x; \mu_2, \sigma_2) + c_3 N(x; \mu_3, \sigma_3) \quad \text{and} \quad \sum c_i = 1 \quad (4)$$

*pdf* parameters:  $c_1, c_2, c_3, \mu_1, \mu_2, \mu_3, \sigma_1, \sigma_2, \sigma_3$

## Observation from GMM

**Experiment:** An urn contains balls of 3 different colours: red, blue or green. Behind a curtain, a person picks a ball from urn

If red ball  $\Rightarrow$  generate  $x[i]$  from  $N(x; \mu_1, \sigma_1)$

If blue ball  $\Rightarrow$  generate  $x[i]$  from  $N(x; \mu_2, \sigma_2)$

If green ball  $\Rightarrow$  generate  $x[i]$  from  $N(x; \mu_3, \sigma_3)$

We have access *only* to observations  $x[0], x[1], \dots, x[N-1]$

Therefore :  $p(x[i]; \theta) = c_1 N(x; \mu_1, \sigma_1) + c_2 N(x; \mu_2, \sigma_2) + c_3 N(x; \mu_3, \sigma_3)$

but we do not know which urn  $x[i]$  comes from!

Can we estimate component  $\theta = [c_1 \ c_2 \ c_3 \ \mu_1 \ \mu_2 \ \mu_3 \ \sigma_1 \ \sigma_2 \ \sigma_3]^T$  from the observations?

$$\arg \max_{\theta} p(\mathbf{X}; \theta) = \arg \max_{\theta} \prod_{i=1}^N p(x_i; \theta) \quad (5)$$

# Estimation of Parameters of GMM

**Easier Problem:** We know the component for each observation

Obs:	x[0]	x[1]	x[2]	x[3]	x[4]	x[5]	x[6]	x[7]	x[8]	x[9]	x[10]	x[11]	x[12]
Comp.	1	2	2	1	3	1	3	3	2	2	3	3	3

$$\mathbf{X}_1 = \{x[0], x[3], x[5]\} \quad \text{belong to} \quad p_1(x; \mu_1, \sigma_1)$$

$$\mathbf{X}_2 = \{x[1], x[2], x[8], x[9]\} \quad \text{belongs to} \quad p_2(x; \mu_2, \sigma_2)$$

$$\mathbf{X}_3 = \{x[3], x[6], x[7], x[10], x[11], x[12]\} \quad \text{belongs to} \quad p_3(x; \mu_3, \sigma_3)$$

From:  $\mathbf{X}_1 = \{x[0], x[3], x[5]\}$

$$\hat{c}_1 = \frac{3}{13}$$

and  $\hat{\mu}_1 = \frac{1}{3} \{x[0] + x[3] + x[5]\}$

$$\hat{\sigma}_1^2 = \frac{1}{3} \{ (x[0] - \hat{\mu}_1)^2 + (x[3] - \hat{\mu}_1)^2 + (x[5] - \hat{\mu}_1)^2 \}$$

In practice we do *not* know which observation come from which *pdf*.

⇒ How do we solve for  $\arg \max_{\theta} p(X; \theta)$  ?



# Incomplete & Complete Data

$x[0], x[1], \dots, x[N - 1] \Rightarrow$  incomplete data,

Introduce another set of variables  $y[0], y[1], \dots, y[N - 1]$

such that  $y[i] = 1$  if  $x[i] \in p_1$ ,  $y[i] = 2$  if  $x[i] \in p_2$  and  $y[i] = 3$  if  $x[i] \in p_3$

Obs:	x[0]	x[1]	x[2]	x[3]	x[4]	x[5]	x[6]	x[7]	x[8]	x[9]	x[10]	x[11]	x[12]
Comp.	1	2	2	1	3	1	3	3	2	2	3	3	3
miss:	y[0]	y[1]	y[2]	y[3]	y[4]	y[5]	y[6]	y[7]	y[8]	y[9]	y[10]	y[11]	y[12]

$y[i] =$  missing data—unobserved data  $\Rightarrow$  information about component

$\mathbf{z} = (\mathbf{x}; \mathbf{y})$  is complete data  $\Rightarrow$  observations and which density they come from

$$p(\mathbf{z}; \boldsymbol{\theta}) = p(\mathbf{x}, \mathbf{y}; \boldsymbol{\theta})$$

$$\Rightarrow \hat{\boldsymbol{\theta}}_{MLE} = \arg \max_{\boldsymbol{\theta}} p(\mathbf{z}; \boldsymbol{\theta})$$

**Question:** But how do we find which observation belongs to which density ?

Given observation  $x[0]$  and  $\theta^g$ , what is the probability of  $x[0]$  coming from first distribution?

$$\begin{aligned}
 & p(y[0] = 1 | x[0]; \theta^g) \\
 &= \frac{p(y[0]=1, x[0]; \theta^g)}{p(x[0]; \theta^g)} \\
 &= \frac{p(x[0] | y[0]=1; \mu_1^g, \sigma_1^g) \cdot p(y[0]=1)}{\sum_{j=1}^3 p(x[0] | y[0]=j; \theta^g) p(y[0]=j)} \\
 &= \frac{p(x[0] | y[0]=1, \mu_1^g, \sigma_1^g) \cdot c_1^g}{p(x[0] | y[0]=1; \mu_1^g, \sigma_1^g) c_1^g + p(x[0] | y[0]=2; \mu_2^g, \sigma_2^g) c_2^g + \dots}
 \end{aligned}$$


All parameters are known  $\Rightarrow$  we can calculate  $p(y[0] = 1 | x[0]; \theta^g)$

(Similarly calculate  $p(y[0] = 2 | x[0]; \theta^g)$ ,  $p(y[0] = 3 | x[0]; \theta^g)$ )

Which density?  $\Rightarrow y[0] = \arg \max_i p(y[0] = i | x[0]; \theta^g)$  – Hard allocation

# Parameter Estimation for Hard Allocation

	$x[0]$	$x[1]$	$x[2]$	$x[3]$	$x[4]$	$x[5]$	$x[6]$
$p(y[j] = 1 x[j]; \theta^g)$	0.5	0.6	0.2	0.1	0.2	0.4	0.2
$p(y[j] = 2 x[j]; \theta^g)$	0.25	0.3	0.75	0.3	0.7	0.5	0.6
$p(y[j] = 3 x[j]; \theta^g)$	0.25	0.1	0.05	0.6	0.1	0.1	0.2
Hard Assign.	$y[0]=1$	$y[1]=1$	$y[2]=2$	$y[3]=3$	$y[4]=2$	$y[5]=2$	$y[6]=2$



Updated Parameters:  $\hat{c}_1 = \frac{2}{7}$     $\hat{c}_2 = \frac{4}{7}$     $\hat{c}_3 = \frac{1}{7}$  (different from initial guess!)

Similarly (for Gaussian) find:  $\hat{\mu}_i, \hat{\sigma}_i^2$  for  $i^{th}$  pdf

## Parameter Estimation for Soft Assignment

	x[0]	x[1]	x[2]	x[3]	x[4]	x[5]	x[6]
$p(y[j] = 1 x[j]; \theta^g)$	0.5	0.6	0.2	0.1	0.2	0.4	0.2
$p(y[j] = 2 x[j]; \theta^g)$	0.25	0.3	0.75	0.3	0.7	0.5	0.6
$p(y[j] = 3 x[j]; \theta^g)$	0.25	0.1	0.05	0.6	0.1	0.1	0.2

Example: Prob. of each sample belonging to component 1

$$p(y[0] = 1|x[0]; \theta^g), p(y[1] = 1|x[1]; \theta^g), p(y[2] = 1|x[2]; \theta^g), \dots$$

Average probability that a sample belongs to Comp.#1 is

$$\begin{aligned}\hat{c}_1^{new} &= \frac{1}{N} \sum_{i=1}^N p(y[i] = 1|x[i]; \theta^g) \\ &= \frac{0.5 + 0.6 + 0.2 + 0.1 + 0.2 + 0.4 + 0.2}{7} = \frac{2.2}{7}\end{aligned}$$

# Soft Assignment – Estimation of Means & Variances

Recall: Prob. of sample  $j$  belonging to component  $i$

$$p(y[j] = i | x[j]; \theta^g)$$

**Soft Assignment:** Parameters estimated by taking weighted average !

$$\mu_1^{new} = \frac{\sum_{i=1}^N x_i \cdot p(y[i] = 1 | x[i]; \theta^g)}{\sum_{i=1}^N p(y[i] = 1 | x[i]; \theta^g)}$$

$$(\sigma_1^2)^{new} = \frac{\sum_{i=1}^N (x_i - \hat{\mu}_1)^2 \cdot p(y[i] = 1 | x[i]; \theta^g)}{\sum_{i=1}^N p(y[i] = 1 | x[i]; \theta^g)}$$

These are updated parameters starting with initial guess  $\theta^g$

# Maximum Likelihood Estimation of Parameters of GMM

1. Make initial guess of parameters:  $\theta^g = c_1^g, c_2^g, c_3^g, \mu_1^g, \mu_2^g, \mu_3^g, \sigma_1^g, \sigma_2^g, \sigma_3^g$
2. Knowing parameters  $\theta^g$ , find Prob. of sample  $x_i$  belonging to  $j^{th}$  component.

$$p[y[i] = j \mid x[i]; \theta^g] \quad \text{for } i = 1, 2, \dots, N \Rightarrow \text{no. of observations}$$
$$\quad \text{for } j = 1, 2, \dots, M \Rightarrow \text{no. of components}$$

3.

$$\hat{c}_j^{new} = \frac{1}{N} \sum_{i=1}^N p(y[i] = j \mid x[i]; \theta^g)$$

4.


$$\mu_j^{new} = \frac{\sum_{i=1}^N x_i \cdot p(y[i] = j \mid x[i]; \theta^g)}{\sum_{i=1}^N p(y[i] = j \mid x[i]; \theta^g)}$$

5.

$$(\sigma_j^2)^{new} = \frac{\sum_{i=1}^N (x_i - \hat{\mu}_1)^2 \cdot p(y[i] = j \mid x[i]; \theta^g)}{\sum_{i=1}^N p(y[i] = j \mid x[i]; \theta^g)}$$

6. Go back to (2) and repeat until convergence

# Gaussian Mixture Example: Start




*Advance apologies: in Black  
and White this example will be  
incomprehensible*

Copyright © 2001, 2004, Andrew W. Moore

Clustering with Gaussian Mixtures: Slide 40


After first  
iteration



Copyright © 2011, 2014, Andrew W. Moore

Clustering with Gaussian Mixtures: Slide 41

After 3rd  
iteration




Copyright © 2011, 2014, Andrew W. Moore

Clustering with Gaussian Mixtures: Slide 43




After 5th iteration



Copyright © 2001, 2004, Andrew W. Moore

Clustering with Gaussian Mixtures: Slide 45

After 20th iteration



Copyright © 2001, 2004, Andrew W. Moore

Clustering with Gaussian Mixtures: Slide 47

Live demonstration: <http://www.neurosci.aist.go.jp/~akaho/MixtureEM.html>

## Practical Issues

- E.M. can get stuck in local minima.
- EM is very sensitive to initial conditions; a good initial guess helps; k-means algorithm is used prior to application of EM algorithm

## Size of a GMM

Bayesian Information Criterion (BIC) value of a GMM can be defined as follows:

$$BIC(G \mid X) = \log p(X \mid \hat{G}) - \frac{d}{2} \log N$$

where

$\hat{G}$  represent the GMM with the ML parameter configuration


$d$  represents the number of parameters in  $G$

$N$  is the size of the dataset

the first term is the log-likelihood term;

the second term is the model complexity penalty term.

BIC selects the best GMM corresponding to the largest BIC value by trading off these two terms.



**Fig. 1.** Data set and the corresponding BIC value curve

source: *Boosting GMM and Its Two Applications*, F.Wang, C.Zhang and N.Lu in N.C.Oza et al. (Eds.) LNCS 3541, pp. 12-21, 2005

The BIC criterion can discover the true GMM size effectively as shown in the figure.

# Maximum A Posteriori (MAP)

- Sometimes, it is difficult to get sufficient number of examples for robust estimation of parameters.
- However, one may have access to large number of similar examples which can be utilized.
- Adapt the target distribution from such a distribution. For example, adapt a speaker independent model to a new speaker using small amount of adaptation data.

# MAP Adaptation

## ML Estimation

$$\hat{\theta}_{MLE} = \arg \max_{\theta} p(X|\theta)$$


## MAP Estimation

$$\begin{aligned}\hat{\theta}_{MAP} &= \arg \max_{\theta} p(\theta|X) \\ &= \arg \max_{\theta} \frac{p(X|\theta) p(\theta)}{p(X)} \\ &= \arg \max_{\theta} p(X|\theta) p(\theta)\end{aligned}$$

$p(\theta)$  is the *a priori* distribution of parameters  $\theta$ .

A conjugate prior is chosen such that the corresponding posterior belongs to the same functional family as the prior.

## Simple Implementation of MAP-GMMs



Source: Statistical Machine Learning from Data: GMM; Samy Bengio

## Simple Implementation

Train a **prior** model **p** with a large amount of available data (say, from multiple speakers). Adapt the parameters to a new speaker using some adaptation data (**X**).

Let  $\alpha = [0, 1]$  be a parameter that describes the faith on the prior model.

Adapted **weight** of  $j^{th}$  mixture

$$\hat{w}_j = \left[ \alpha w_j^p + (1 - \alpha) \sum_i p(j|x_i) \right] \gamma$$

Here  $\gamma$  is a normalization factor such that  $\sum w_j = 1$ .

## Simple Implementation (contd.)

means

$$\hat{\mu}_j = \alpha \mu_j^p + (1 - \alpha) \frac{\sum_i p(j|x_i) x_i}{\sum_i p(j|x_i)}$$

Weighted average of sample mean and *a prior* mean

variances

$$\hat{\sigma}_j = \alpha \left( \sigma_j^p + \mu_j^p \mu_j^{p'} \right) + (1 - \alpha) \frac{\sum_i p(j|x_i) x_i x_i'}{\sum_i p(j|x_i)} - \hat{\mu}_j \hat{\mu}_j'$$



# HMM

- Primary role of speech signal is to carry a message; sequence of sounds (**phonemes**) encode a sequence of words.
- The acoustic manifestation of a phoneme is mostly determined by:
  - Configuration of articulators (jaw, tongue, lip)
  - physiology and emotional state of speaker
  - Phonetic context
- HMM models sequential patterns; speech is a sequential pattern
- Most text dependent speaker recognition systems use HMMs
- Text verification involves verification/recognition of phonemes

# Phoneme recognition

Consider two phonemes classes /aa/ and /iy/.

**Problem:** Determine to which class a given sound belongs.

Processing of speech signal results in a sequence of feature (observation) vectors:  
 $\mathbf{o}_1, \dots, \mathbf{o}_T$  (say MFCC vectors)

We say the speech is /aa/ if:  $p(aa|\mathbf{O}) > p(iy|\mathbf{O})$

Using Bayes Rule

$$\frac{\overbrace{p(\mathbf{O}|aa)}^{AcousticModel} p(aa)}{p(\mathbf{O})} \quad V.s. \quad \frac{p(\mathbf{O}|iy) \overbrace{p(iy)}^{PriorProb}}{p(\mathbf{O})}$$

Given  $p(\mathbf{O}|aa)$ ,  $p(aa)$ ,  $p(\mathbf{O}|iy)$  and  $p(iy)$   $\Rightarrow$  which is more probable ?

# Parameter Estimation of Acoustic Model

How do we find the density function  $p_{aa}(\cdot)$  and  $p_{iy}(\cdot)$ .

We assume a parametric model:  $\Rightarrow p_{aa}()$  parameterised by  $\theta_{aa}$   
 $\Rightarrow p_{ij}()$  parameterised by  $\theta_{iy}$

Training Phase: Collect many examples of /aa/ being said  
 $\Rightarrow$  Compute corresponding observations  $\mathbf{o}_1, \dots, \mathbf{o}_{T_{aa}}$

Use the *Maximum Likelihood Principle*

$$\widehat{\theta_{aa}} = \arg \max_{\theta_{aa}} p(\mathbf{O}; \theta_{aa})$$

Recall: if the *pdf* is modelled as a Gaussian Mixture Model  
 $\Rightarrow$  then we use EM Algorithm

# Modelling of Phoneme

Our Articulators are moving from a configuration  
To enunciate /aa/ in a word  $\Rightarrow$  for previous phoneme to /aa/ and then proceeding  
to move to configuration of next phoneme.

Can think of 3 distinct time periods:






- $\Rightarrow$  Transition from previous phoneme
- $\Rightarrow$  Steady state
- $\Rightarrow$  Transition to next phoneme

Features for 3 “time-interval ” are quite different

- $\Rightarrow$  Use different density functions to model the three time intervals
- $\Rightarrow$  model as  $p_{aa^1}(\cdot; \theta_{aa^1}) \quad p_{aa^2}(\cdot; \theta_{aa^2}) \quad p_{aa^3}(\cdot; \theta_{aa^3})$


Also need to model the *time durations* of these time-intervals – transition probs.

# Stochastic Model (HMM)



# HMM Model of Phoneme

- Use term “State” for each of the three time periods.
- Prob. of  $\mathbf{o}_t$  from  $j^{th}$  state, i.e.  $p_{aa^j}(\mathbf{o}_t; \boldsymbol{\theta}_{aa^j}) \Rightarrow$  denoted as  $b_j(\mathbf{o}_t)$




- Observation,  $\mathbf{o}_t$ , is generated by which state density?
  - Only observations are seen, the state-sequence is “hidden”
  - Recall: In GMM, the “mixture component is “hidden”

## Probability of Observation

Recall: To classify, we evaluate  $Pr(\mathbf{O}|\Lambda)$  – where  $\Lambda$  are parameters of models

In /aa/ Vs /iy/ calculation:  $\Rightarrow Pr(\mathbf{O}|\Lambda_{aa})$  Vs  $Pr(\mathbf{O}|\Lambda_{iy})$

**Example:** 2-state HMM model and 3 observations  $\mathbf{o}_1 \mathbf{o}_2 \mathbf{o}_3$




Model parameters are assumed known:

Transition Prob.  $\Rightarrow a_{11}$ ,  $a_{12}$ ,  $a_{21}$ , and  $a_{22}$  – model time durations

State density  $\Rightarrow b_1(\mathbf{o}_t)$  and  $b_2(\mathbf{o}_t)$ .

$b_j(\mathbf{o}_t)$  are usually modelled as single Gaussian with parameter  $\mu_j$ ,  $\sigma_j^2$  or by GMMs

# Probability of Observation through one Path



$T = 3$  observations and  $N = 2$  nodes  $\Rightarrow 8$  paths thru 2 nodes for 3 observations

**Example:** Path  $P_1$  through states 1, 1, 1.

$$Pr\{\mathbf{O}|P_1, \Lambda\} = b_1(\mathbf{o}_1) \cdot b_1(\mathbf{o}_2) \cdot b_1(\mathbf{o}_3)$$

$$\text{Prob. of Path } P_1 = Pr\{P_1|\Lambda\} = a_{01} \cdot a_{11} \cdot a_{11}$$

$$Pr\{\mathbf{O}, P_1|\Lambda\} = Pr\{\mathbf{O}|P_1, \Lambda\} \cdot Pr\{P_1|\Lambda\} = a_{01}b_1(\mathbf{o}_1).a_{11}b_1(\mathbf{o}_2).a_{11}b_1(\mathbf{o}_3)$$



## Probability of Observation

Path	$\mathbf{o}_1$	$\mathbf{o}_2$	$\mathbf{o}_3$	$p(\mathbf{O}, P_i \Lambda)$
$P_1$	1	1	1	$a_{01}b_1(\mathbf{o}_1).a_{11}b_1(\mathbf{o}_2).a_{11}b_1(\mathbf{o}_3)$
$P_2$	1	1	2	$a_{01}b_1(\mathbf{o}_1).a_{11}b_1(\mathbf{o}_2).a_{12}b_2(\mathbf{o}_3)$
$P_3$	1	2	1	$a_{01}b_1(\mathbf{o}_1).a_{12}b_2(\mathbf{o}_2).a_{21}b_1(\mathbf{o}_3)$
$P_4$	1	2	2	$a_{01}b_1(\mathbf{o}_1).a_{12}b_2(\mathbf{o}_2).a_{22}b_2(\mathbf{o}_3)$
$P_5$	2	1	1	$a_{02}b_2(\mathbf{o}_1).a_{21}b_1(\mathbf{o}_2).a_{11}b_1(\mathbf{o}_3)$
$P_6$	1	1	2	$a_{02}b_2(\mathbf{o}_1).a_{21}b_1(\mathbf{o}_2).a_{12}b_2(\mathbf{o}_3)$
$P_7$	1	1	1	$a_{02}b_2(\mathbf{o}_1).a_{22}b_1(\mathbf{o}_2).a_{21}b_1(\mathbf{o}_3)$
$P_8$	1	1	2	$a_{02}b_2(\mathbf{o}_1).a_{22}b_1(\mathbf{o}_2).a_{22}b_2(\mathbf{o}_3)$


$$p(\mathbf{O}|\Lambda) = \sum_{P_i} P\{\mathbf{O}, P_i|\Lambda\} = \sum_{P_i} P\{\mathbf{O}|P_i, \Lambda\} \cdot P\{P_i, \Lambda\}$$

Forward Algorithm  $\Rightarrow$  Avoid Repeat Calculations:

Two Multiplications

$$\begin{aligned} & \overbrace{a_{01}b_1(\mathbf{o}_1)a_{11}b_1(\mathbf{o}_2).a_{11}b_1(\mathbf{o}_3) + a_{02}b_2(\mathbf{o}_1).a_{21}b_1(\mathbf{o}_2).a_{11}b_1(\mathbf{o}_3)} \\ &= \underbrace{[a_{01}.b_1(\mathbf{o}_1)a_{11}b_1(\mathbf{o}_2) + a_{02}b_2(\mathbf{o}_1)a_{21}b_1(\mathbf{o}_2)]a_{11}b_1(\mathbf{o}_3)}_{\text{One Multiplication}} \end{aligned}$$

## Forward Algorithm – Recursion



$$\text{Let } \alpha_1(t=1) = a_{01}b_1(\mathbf{o}_1)$$

$$\text{Let } \alpha_2(t=1) = a_{02}b_2(\mathbf{o}_2)$$

$$\begin{aligned} \text{Recursion : } \alpha_1(t=2) &= [a_{01}b_1(\mathbf{o}_1).a_{11} + a_{02}b_2(\mathbf{o}_1).a_{21}].b_1(\mathbf{o}_2) \\ &= [\alpha_1(t=1).a_{11} + \alpha_2(t=1).a_{21}].b_1(\mathbf{o}_2) \end{aligned}$$

## General Recursion in Forward Algorithm

$$\begin{aligned}\alpha_j(t) &= \left[ \sum \alpha_i(t-1) a_{ij} \right] \cdot b_j(\mathbf{o}_t) \\ &= P\{\mathbf{o}_1, \mathbf{o}_2, \dots, \mathbf{o}_t, s_t = j | \Lambda\}\end{aligned}$$

### Note

$\alpha_j(t) \Rightarrow$  Sum of probabilities of all paths ending at node  $j$  at time  $t$  with partial observation sequence  $\mathbf{o}_1, \mathbf{o}_2, \dots, \mathbf{o}_t$

The probability of the entire observation  $(\mathbf{o}_1, \mathbf{o}_2, \dots, \mathbf{o}_T)$ , therefore, is


$$\begin{aligned}p(\mathbf{O} | \Lambda) &= \sum_{j=1}^N P\{\mathbf{o}_1, \mathbf{o}_2, \dots, \mathbf{o}_T, S_T = j | \Lambda\} \\ &= \sum_{j=1}^N \alpha_j(T)\end{aligned}$$

where  $N = \text{No. of nodes}$

## Backward Algorithm

- analogous to Forward, but coming from the last time instant T

Example:  $a_{01}b_1(\mathbf{o}_1).a_{11}b_1(\mathbf{o}_2).a_{11}b_1(\mathbf{o}_3) + a_{01}b_1(\mathbf{o}_1).a_{11}b_1(\mathbf{o}_2)a_{12}b_2(\mathbf{o}_3) + \dots$   
 $= [a_{01}.b_1(\mathbf{o}_1).a_{11}b_1(\mathbf{o}_2)].(a_{11}b_1(\mathbf{o}_3) + a_{12}b_2(\mathbf{o}_3))$



$$\begin{aligned}
 \beta_1(t=2) &= p\{\mathbf{o}_3 | s_{t=2} = 1, \Lambda\} \\
 &= p\{\mathbf{o}_3, s_{t=3} = 1 | s_{t=2} = 1; \Lambda\} + p\{\mathbf{o}_3, s_{t=3} = 2 | s_{t=2} = 1, \Lambda\} \\
 &= p\{\mathbf{o}_3 | s_{t=3} = 1, s_{t=2} = 1, \Lambda\}.p\{s_{t=3} = 1 | s_{t=2} = 1, \Lambda\} + \\
 &= p\{\mathbf{o}_3 | s_{t=3} = 2, s_{t=2} = 1, \Lambda\}.p\{s_{t=3} = 2 | s_{t=2} = 1, \Lambda\} \\
 &= b_1(\mathbf{o}_3).a_{11} + b_2(\mathbf{o}_3).a_{12}
 \end{aligned}$$

## General Recursion in Backward Algorithm

Given that we are at node  $j$  at time  $t$   
 $\beta_j(t) \Rightarrow$  Sum of probabilities of all paths such that  
partial sequence  $\mathbf{o}_{t+1}, \dots, \mathbf{o}_T$  are observed

$$\beta_i(t) = \underbrace{\sum_{j=1}^N [a_{ij} b_j(\mathbf{o}_{t+1})]}_{\text{Going to each node from } i^{th} \text{ node}}$$

$\underbrace{\beta_j(t+1)}_{\text{Prob. of observation } \mathbf{o}_{t+2} \dots \mathbf{o}_T \text{ given}} \\ \text{now we are in } j^{th} \text{ node at } t+1$

$$= p\{\mathbf{o}_{t+1}, \dots, \mathbf{o}_T | s_t = i, \Lambda\}$$

# Estimation of Parameters of HMM Model

- Given known Model parameters,  $\Lambda$ :
  - Evaluated  $p(\mathbf{O}|\Lambda) \Rightarrow$  useful for classification
  - Efficient Implementation: Use Forward or Backward Algo.
- Given set of observation vectors,  $\mathbf{o}_t$  how do we estimate parameters of HMM?
  - Do not know which states  $\mathbf{o}_t$  come from
    - \* Analogous to GMM – do not know which component
  - Use a special case of EM – Baum-Welch Algorithm
  - Use following relations from Forward/Backward

$$p(\mathbf{O}|\Lambda) = \alpha_N(T) = \beta_1(T) = \sum_{j=1}^N \alpha_j(t) \beta_j(t)$$

# Parameter Estimation for Known State Sequence

Assume each state is modelled as a single Gaussian:

$\hat{\mu}_j$  = Sample mean of observations assigned to state  $j$ .

$\hat{\sigma}_j^2$  = Variance of the observations assigned to state  $j$ .

and

Trans. Prob. from state  $i$  to  $j$  =  $\frac{\text{No. of times transition was made from } i \text{ to } j}{\text{Total number of times we made transition from } i}$

In practice since we do not know which state generated the observation

⇒ So we will do probabilistic assignment.

# Review of GMM Parameter Estimation

Do not know which component of the GMM generated output observation.

Given initial model parameters  $\Lambda^g$ , and observation sequence  $x_1, \dots, x_T$ .

Find probability  $x_i$  comes from component  $j \Rightarrow$  Soft Assignment

$$p[\text{component} = 1 | x_i; \Lambda^g] = \frac{p[\text{component} = 1, x_i | \Lambda_g]}{p[x_i | \Lambda_g]}$$

So, re-estimation equations are:

$$\begin{aligned}\hat{C}_j &= \frac{1}{T} \sum_{i=1}^T p(\text{comp} = j | x_i; \Lambda_g) \\ \hat{\mu}_j^{new} &= \frac{\sum_{i=1}^T x_i p(\text{comp} = j | x_i; \Lambda^g)}{\sum_{i=1}^T p(\text{comp} = j | x_i; \Lambda^g)} \quad \hat{\sigma}_j^2 = \frac{\sum_{i=1}^T (x_i - \hat{\mu}_j)^2 p(\text{comp} = j | x_i; \Lambda_g)}{\sum_{i=1}^T p(\text{comp} = j | x_i; \Lambda^g)}\end{aligned}$$

A similar analogy holds for hidden Markov models



# Baum-Welch Algorithm

Here: We do not know which observation  $\mathbf{o}_t$  comes from which state  $s_i$

Again like GMM we will assume initial guess parameter  $\Lambda^g$

Then prob. of being in “state= $i$  at time= $t$ ” and “state= $j$  at time= $t+1$ ” is

$$\hat{\tau}_t(i, j) = p\{q_t = i, q_{t+1} = j | \mathbf{O}, \Lambda^g\} = \frac{p\{q_t = i, q_{t+1} = j, \mathbf{O} | \Lambda^g\}}{p\{\mathbf{O} | \Lambda^g\}}$$

where  $p\{\mathbf{O} | \Lambda^g\} = \alpha_N(T) = \sum_i \alpha_i(T)$

## Baum-Welch Algorithm

Then prob. of being in “state= $i$  at time= $t$ ” and “state= $j$  at time= $t+1$ ” is

$$\hat{\tau}_t(i, j) = p\{q_t = i, q_{t+1} = j | \mathbf{O}, \Lambda^g\} = \frac{p\{q_t = i, q_{t+1} = j, \mathbf{O} | \Lambda^g\}}{p\{\mathbf{O} | \Lambda^g\}}$$

where  $p\{\mathbf{O} | \Lambda^g\} = \alpha_N(T) = \sum_i \alpha_i(T)$

From ideas of Forward-Backward Algorithm, numerator is

$$p\{q_t = i, q_{t+1} = j, \mathbf{O} | \Lambda^g\} = \alpha_i(t) \cdot a_{ij} b_j(\mathbf{o}_{t+1}) \cdot \beta_j(t+1)$$

$$\text{So } \hat{\tau}_t(i, j) = \frac{\alpha_i(t) \cdot a_{ij} b_j(\mathbf{o}_{t+1}) \beta_j(t+1)}{\alpha_N(t)}$$

## Estimating Transition Probability

Trans. Prob. from state  $i$  to  $j = \frac{\text{No. of times transition was made from } i \text{ to } j}{\text{Total number of times we made transition from } i}$

$\hat{\tau}_t(i, j) \Rightarrow$  prob. of being in “state= $i$  at time= $t$ ” and “state= $j$  at time= $t+1$ ”

If we average  $\hat{\tau}_t(i, j)$  over all time-instants, we get the number of times the system was in  $i^{th}$  state and made a transition to  $j^{th}$  state. So, a revised estimation of transition probability is

$$\hat{a}_{ij}^{new} = \frac{\sum_{t=1}^{T-1} \tau_t(i, j)}{\sum_{t=1}^T \left( \underbrace{\sum_{j=1}^N \tau_t(i, j)}_{\text{all transitions out of } i \text{ at time}=t} \right)}$$

## Estimating State-Density Parameters

Analogous to GMM: which observation belonged to which component,

New estimates for the state *pdf* parameters are (assuming single Gaussian)

$$\hat{\mu}_i = \frac{\sum_{t=1}^T \gamma_i(t) \mathbf{o}_t}{\sum_{t=1}^T \gamma_i(t)}$$

$$\hat{\Sigma}_i = \frac{\sum_{t=1}^T \gamma_i(t) (\mathbf{o}_t - \hat{\mu}_i)(\mathbf{o}_t - \hat{\mu}_i)^T}{\sum_{t=1}^T \gamma_i(t)}$$

These are weighted averages  $\Rightarrow$  weighted by Prob. of being in state  $j$  at  $t$

- Given observation  $\Rightarrow$  HMM model parameters estimated iteratively
- $p(\mathbf{O}|\Lambda) \Rightarrow$  evaluated efficiently by Forward/Backward algorithm

# Viterbi Algorithm

Given the observation sequence,

- the goal is to find corresponding state-sequence that generated it
- there are many possible combination ( $N^T$ ) of state sequence  $\Rightarrow$  many paths.

One possible criterion : Choose the state sequence corresponding to path that with maximum probability

$$\max_i P\{\mathbf{O}, P_i | \Lambda\}$$

Word : represented as sequence of phones

Phone : represented as sequence states

Optimal state-sequence  $\Rightarrow$  Optimal phone-sequence  $\Rightarrow$  Word sequence

# Viterbi Algorithm and Forward Algorithm

Recall Forward Algorithm : We found probability over each path and summed over all possible paths


$$\sum_{i=1}^{N^T} p\{\mathbf{O}, P_i | \Lambda\}$$

Viterbi is just special case of Forward algo.

At each node  $\left\{ \begin{array}{l} \text{instead of sum of prob. of all paths} \\ \text{choose path with max prob.} \end{array} \right.$

In Practice:  $p(\mathbf{O} | \Lambda)$  approximated by Viterbi (instead of Forward Algo.)

# Viterbi Algorithm



# Decoding


- Recall: Desired transcription  $\widehat{\mathbf{W}}$  obtained by maximising

$$\widehat{\mathbf{W}} = \arg \max_{\mathbf{W}} p(\mathbf{W}|\mathbf{O})$$

- Search over all possible  $\mathbf{W}$  – astronomically large!
- Viterbi Search – find most likely path through a HMM
  - Sequence of phones (states) which is most probable
  - Mostly: most probable sequence of phones correspond to most probable sequence of words

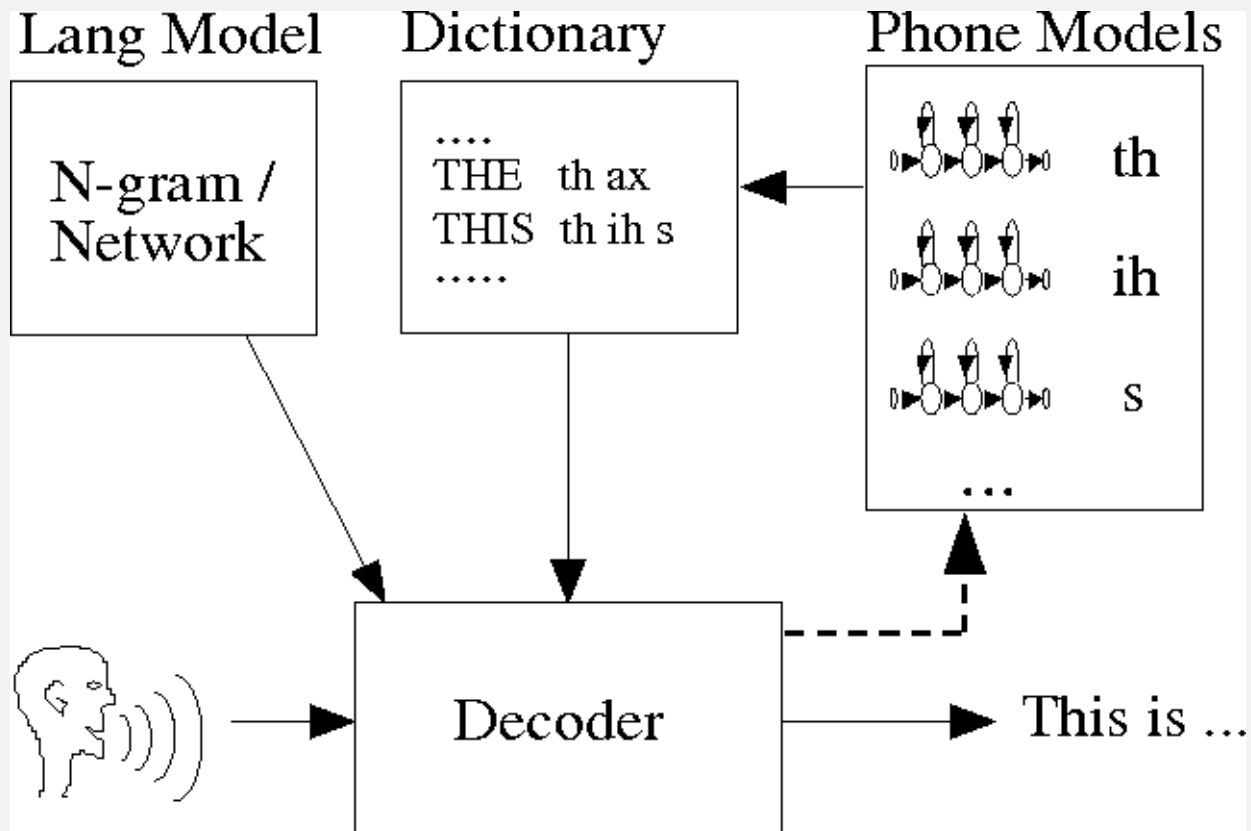


# Training of a Speech Recognition System




- HMM parameter's estimated using large databases – 100 hours
  - \* Parameters estimated using Maximum Likelihood Criterion

# Recognition



# Speaker Recognition

Spectra (formants) of a given sound are different for different speakers.



Spectra of 2 speakers for *one* “frame” of /iy/

Derive speaker dependent model of a new speaker by MAP adaptation of Speaker-Independent (SI) model using small amount of adaptation data; use for speaker recognition.

## References

- *Pattern Classification*, R.O.Duda, P.E.Hart and D.G.Stork, John Wiley, 2001.
- *Introduction to Statistical Pattern Recognition*, K.Fukunaga, Academic Press, 1990.
- *The EM Algorithm and Extensions*, Geoffrey J. McLachlan and Thriyambakam Krishnan, Wiley-Interscience; 2nd edition, 2008. ISBN-10: 0471201707
- *The EM Algorithm and Extensions*, Geoffrey J. McLachlan and Thriyambakam Krishnan, Wiley-Interscience; 2nd edition, 2008. ISBN-10: 0471201707
- *Fundamentals of Speech Recognition*, Lawrence Rabiner & Biing-Hwang Juang, Englewood Cliffs NJ: PTR Prentice Hall (Signal Processing Series), c1993, ISBN 0-13-015157-2

- *Hidden Markov models for speech recognition*, X.D. Huang, Y. Ariki, M.A. Jack. Edinburgh: Edinburgh University Press, c1990.
- *Statistical methods for speech recognition*, F.Jelinek, The MIT Press, Cambridge, MA., 1998.
- *Maximum Likelihood from incomplete data via the em algorithm*, J. Royal Statistical Soc. **39**(1), pp. 1-38, 1977.
- Maximum a *Posteriori* Estimation for Multivariate Gaussian Mixture Observation of Markov Chains, J.-L.Gauvain and C.-H.Lee, IEEE Trans. SAP, **2**(2), pp. 291-298, 1994.
- *A Gentle Tutorial of the EM Algorithm and its Application to Parameter Estimation for Gaussian Mixture and Hidden Markov Models*, J.A.Bilmes, ISCI, TR-97-021.
- *Boosting GMM and Its Two Applications*, F.Wang, C.Zhang and N.Lu in N.C.Oza et al. (Eds.) LNCS 3541, pp. 12-21, 2005.