

CS 224S / LINGUIST 281

Speech Recognition, Synthesis, and Dialogue

Dan Jurafsky

Lecture 4: Intro to Festival; rest of Text Normalization; Letter-to-Sound

IP Notice: lots of info, text, and diagrams on these slides comes (thanks!) from Alan Black's excellent lecture notes and from Richard Sproat's great new slides.

Outline

- Overview of Festival
 - ♦ Where it lives, its components
 - ♦ Its scripting language: Scheme
- Finishing up Part of Speech Tagging
- Phonetic Analysis
 - ♦ Dictionaries
 - ♦ Names
 - ♦ Letter-to-Sound Rules
 - (or “Grapheme-to-Phoneme Conversion”)

Festival

- Open source speech synthesis system
- Designed for development and runtime use
 - ◆ Use in many commercial and academic systems
 - ◆ Distributed with RedHat 9.x, etc
 - ◆ Hundreds of thousands of users
 - ◆ Multilingual
 - No built-in language
 - Designed to allow addition of new languages
 - ◆ Additional tools for rapid voice development
 - Statistical learning tools
 - Scripts for building models

Festival as software

- <http://festvox.org/festival/>
- General system for multi-lingual TTS
- C/C++ code with Scheme scripting language
- General replaceable modules:
 - ♦ Lexicons, LTS, duration, intonation, phrasing, POS tagging, tokenizing, diphone/unit selection, signal processing
- General tools
 - ♦ Intonation analysis (f0, Tilt), signal processing, CART building, N-gram, SCFG, WFST

Festival as software

- <http://festvox.org/festival/>
- No fixed theories
- New languages without new C++ code
- Multiplatform (Unix/Windows)
- Full sources in distribution
- Free software

CMU FestVox project

- Festival is an engine, how do you make voices?
- Festvox: building synthetic voices:
 - ♦ Tools, scripts, documentation
 - ♦ Discussion and examples for building voices
 - ♦ Example voice databases
 - ♦ Step by step walkthroughs of processes
- Support for English and other languages
- Support for different waveform synthesis methods
 - ♦ Diphone
 - ♦ Unit selection
 - ♦ Limited domain

Synthesis tools

- I want my computer to talk
 - ♦ Festival Speech Synthesis
- I want my computer to talk in my voice
 - ♦ FestVox Project
- I want it to be fast and efficient
 - ♦ Flite

Using Festival

- How to get Festival to talk
- Scheme (Festival's scripting language)
- Basic Festival commands

Getting it to talk

- Say a file
 - ♦ `festival --tts file.txt`
- From Emacs
 - ♦ `say region, say buffer`
- Command line interpreter
 - ♦ `festival> (SayText "hello")`

Scheme: the scripting lg

- Advantages of a scripting lg
 - ◆ Convenient, easy to add functionality
- Why Scheme?
 - ◆ Holdover from the LISP days of AI.
 - ◆ Many people like it.
 - ◆ It's very simple

Quick Intro to Scheme

- Scheme is a dialect of LISP
- **expressions** are
 - ♦ **atoms** or
 - ♦ **lists**
 - a bcd "hello world" 12.3**
 - (a b c)**
 - (a (1 2) seven)**
- Interpreter evaluates expressions
 - ♦ Atoms evaluate as variables
 - ♦ Lists evaluate as functional calls
 - bx x**
 - 3.14**
 - (+ 2 3)**

Quick Intro to Scheme

- Setting variables

```
(set! a 3.14)
```

- Defining functions

```
(define (timestwo n) (* 2 n))
```

```
(timestwo a)
```

```
6.28
```

Lists in Scheme

- **festival> (set! alist '(apples pears bananas))**
- (apples pears bananas)
- **festival> (car alist)**
- apples
- **festival> (cdr alist)**
- (pears bananas)
- **festival> (set! blist (cons 'oranges alist))**
- (oranges apples pears bananas)
- **festival> append alist blist**
- #<SUBR(6) append>
- (apples pears bananas)
- (oranges apples pears bananas)
- **festival> (append alist blist)**
- (apples pears bananas oranges apples pears bananas)
- **festival> (length alist)**
- 3
- **festival> (length (append alist blist))**
- 7

Scheme: speech

- Make an utterance of type text

```
festival> (set! utt1 (Utterance Text "hello"))
#<Utterance 0xf6855718>
```
- Synthesize an utterance

```
festival> (utt.synth utt1)
#<Utterance 0xf6855718>
```
- Play waveform

```
festival> (utt.play utt1)
#<Utterance 0xf6855718>
```
- Do all together

```
festival> (SayText "This is an example")
#<Utterance 0xf6961618>
```

Scheme: speech

- In a file

```
(define (SpeechPlus a b)
  (SayText
    (format nil
      "%d plus %d equals %d"
      a b (+ a b))))
```
- Loading files

```
festival> (load "file.scm")
t
```
- Do all together

```
festival> (SpeechPlus 2 4)
#<Utterance 0xf6961618>
```

Scheme: speech

```
(define (sp_time hour minute)
  (cond
    (( < hour 12)
     (SayText
      (format nil
        "It is %d %d in the morning"
        hour minute )))
    (( < hour 18)
     (SayText
      (format nil
        "It is %d %d in the afternoon"
        (- hour 12) minute )))
    (t
     (SayText
      (format nil
        "It is %d %d in the evening"
        (- hour 12) minute )))))
```


Getting help

- Online manual
 - ♦ <http://festvox.org/docs/manual-1.4.3>
- Alt-h (or esc-h) on current symbol short help
- Alt-s (or esc-s) to speak help
- Alt-m goto man page
- Use TAB key for completion

Festival Structure

Utterance structure in Festival

- http://www.festvox.org/docs/manual-1.4.2/festival_14.html
- Features in festival
- http://www.festvox.org/docs/manual-1.4.2/festival_32.html

Hidden Markov Model Tagging

- Using an HMM to do POS tagging
- Is a special case of Bayesian inference
 - ◆ Foundational work in computational linguistics
 - ◆ Bledsoe 1959: OCR
 - ◆ Mosteller and Wallace 1964: authorship identification
- It is also related to the “noisy channel” model that we’ll do when we do ASR (speech recognition)

POS tagging as a sequence classification task

- We are given a sentence (an “observation” or “sequence of observations”)
 - ♦ Secretariat is expected to race tomorrow
- What is the best sequence of tags which corresponds to this sequence of observations?
- Probabilistic view:
 - ♦ Consider all possible sequences of tags
 - ♦ Out of this universe of sequences, choose the tag sequence which is most probable given the observation sequence of n words $w_1 \dots w_n$.

Getting to HMM

- We want, out of all sequences of n tags $t_1 \dots t_n$ the single tag sequence such that $P(t_1 \dots t_n | w_1 \dots w_n)$ is highest.

$$\hat{t}_1^n = \operatorname{argmax}_{t_1^n} P(t_1^n | w_1^n)$$

- Hat $\hat{}$ means “our estimate of the best one”
- $\operatorname{Argmax}_x f(x)$ means “the x such that $f(x)$ is maximized”

Getting to HMM

- This equation is guaranteed to give us the best tag sequence

$$\hat{t}_1^n = \operatorname{argmax}_{t_1^n} P(t_1^n | w_1^n)$$

- But how to make it operational? How to compute this value?
- Intuition of Bayesian classification:
 - ♦ Use Bayes rule to transform into a set of other probabilities that are easier to compute

Using Bayes Rule

$$P(x|y) = \frac{P(y|x)P(x)}{P(y)}$$

$$\hat{t}_1^n = \operatorname{argmax}_{t_1^n} \frac{P(w_1^n | t_1^n) P(t_1^n)}{P(w_1^n)}$$

$$\hat{t}_1^n = \operatorname{argmax}_{t_1^n} P(w_1^n | t_1^n) P(t_1^n)$$

Likelihood and prior

$$\hat{t}_1^n = \operatorname{argmax}_{t_1^n} \overbrace{P(w_1^n | t_1^n)}^{\text{likelihood}} \overbrace{P(t_1^n)}^{\text{prior}}$$

$$P(w_1^n | t_1^n) \approx \prod_{i=1}^n P(w_i | t_i)$$

$$P(t_1^n) \approx \prod_{i=1}^n P(t_i | t_{i-1})$$

$$\hat{t}_1^n = \operatorname{argmax}_{t_1^n} P(t_1^n | w_1^n) \approx \operatorname{argmax}_{t_1^n} \prod_{i=1}^n P(w_i | t_i) P(t_i | t_{i-1})$$

Two kinds of probabilities (1)

- Tag transition probabilities $p(t_i|t_{i-1})$
 - ♦ Determiners likely to precede adjs and nouns
 - That/DT flight/NN
 - The/DT yellow/JJ hat/NN
 - So we expect $P(NN|DT)$ and $P(JJ|DT)$ to be high
 - But $P(DT|JJ)$ to be:
 - ♦ Compute $P(NN|DT)$ by counting in a labeled corpus:

$$P(t_i|t_{i-1}) = \frac{C(t_{i-1}, t_i)}{C(t_{i-1})}$$

$$P(NN|DT) = \frac{C(DT, NN)}{C(DT)} = \frac{56,509}{116,454} = .49$$

Two kinds of probabilities (2)

- Word likelihood probabilities $p(w_i|t_i)$
 - ♦ VBZ (3sg Pres verb) likely to be “is”
 - ♦ Compute $P(\text{is}|\text{VBZ})$ by counting in a labeled corpus:

$$P(w_i|t_i) = \frac{C(t_i, w_i)}{C(t_i)}$$

$$P(\text{is}|\text{VBZ}) = \frac{C(\text{VBZ}, \text{is})}{C(\text{VBZ})} = \frac{10,073}{21,627} = .47$$

POS tagging: likelihood and prior

$$\hat{t}_1^n = \operatorname{argmax}_{t_1^n} \overbrace{P(w_1^n | t_1^n)}^{\text{likelihood}} \overbrace{P(t_1^n)}^{\text{prior}}$$

$$P(w_1^n | t_1^n) \approx \prod_{i=1}^n P(w_i | t_i)$$



$$P(t_1^n) \approx \prod_{i=1}^n P(t_i | t_{i-1})$$

$$\hat{t}_1^n = \operatorname{argmax}_{t_1^n} P(t_1^n | w_1^n) \approx \operatorname{argmax}_{t_1^n} \prod_{i=1}^n P(w_i | t_i) P(t_i | t_{i-1})$$

An Example: the verb “race”

- Secretariat/**NNP** is/**VBZ** expected/**VCN** to/**TO** **race**/**VB** tomorrow/**NR**
- People/**NNS** continue/**VB** to/**TO** inquire/**VB** the/**DT** reason/**NN** for/**IN** the/**DT** **race**/**NN** for/**IN** outer/**JJ** space/**NN**
- How do we pick the right tag?

Disambiguating “race”



- $P(\text{NN}|\text{TO}) = .00047$
- $P(\text{VB}|\text{TO}) = .83$
- $P(\text{race}|\text{NN}) = .00057$
- $P(\text{race}|\text{VB}) = .00012$
- $P(\text{NR}|\text{VB}) = .0027$
- $P(\text{NR}|\text{NN}) = .0012$

- $P(\text{VB}|\text{TO})P(\text{NR}|\text{VB})P(\text{race}|\text{VB}) = .00000027$
- $P(\text{NN}|\text{TO})P(\text{NR}|\text{NN})P(\text{race}|\text{NN}) = .00000000032$

- So we (correctly) choose the verb reading

Hidden Markov Models

- What we've described with these two kinds of probabilities is a Hidden Markov Model
- A Hidden Markov Model is a particular probabilistic kind of automaton
- Let's just spend a bit of time tying this into the model
- We'll return to this in much more detail in 2 weeks when we do ASR

Hidden Markov Model

$$Q = q_1 q_2 \dots q_N$$

a set of N **states**.

$$A = a_{11} a_{12} \dots a_{n1} \dots a_{nn}$$

a **transition probability matrix** A , each a_{ij} representing the probability of moving from state i to state j , s.t. $\sum_{j=1}^n a_{ij} = 1 \quad \forall i$.

$$O = o_1 o_2 \dots o_T$$

a sequence of T **observations**, each one drawn from a vocabulary $V = v_1, v_2, \dots, v_V$.


$$B = b_i(o_t)$$

A sequence of **observation likelihoods**, also called **emission probabilities**, each expressing the probability of an observation o_t being generated from a state i .


$$q_0, q_F$$

a special **start state** and **end (final) state** that are not associated with observations, together with transition probabilities $a_{01} a_{02} \dots a_{0n}$ out of the start state and $a_{1F} a_{2F} \dots a_{nF}$ into the end state.

Transitions between the hidden states of HMM, showing A probs



B observation likelihoods for POS HMM



The A matrix for the POS HMM




	VB	TO	NN	PPSS
<s>	.019	.0043	.041	.067
VB	.0038	.035	.047	.0070
TO	.83	0	.00047	0
NN	.0040	.016	.087	.0045
PPSS	.23	.00079	.0012	.00014

The B matrix for the POS HMM



	I	want	to	race
VB	0	.0093	0	.00012
TO	0	0	.99	0
NN	0	.000054	0	.00057
PPSS	.37	0	0	0

Viterbi intuition: we are looking for the best 'path'



The Viterbi Algorithm

function VITERBI(*observations* of len T , *state-graph* of len N) **returns** *best-path*

create a path probability matrix $viterbi[N+2, T]$

for each state s **from** 1 **to** N **do** ; initialization step

$viterbi[s, 1] \leftarrow a_{0,s} * b_s(o_1)$

$backpointer[s, 1] \leftarrow 0$

for each time step t **from** 2 **to** T **do** ; recursion step

for each state s **from** 1 **to** N **do**

$viterbi[s, t] \leftarrow \max_{s'=1}^N viterbi[s', t-1] * a_{s',s} * b_s(o_t)$

$backpointer[s, t] \leftarrow \operatorname{argmax}_{s'=1}^N viterbi[s', t-1] * a_{s',s}$

$viterbi[q_F, T] \leftarrow \max_{s=1}^N viterbi[s, T] * a_{s,q_F}$; termination step

$backpointer[q_F, T] \leftarrow \operatorname{argmax}_{s=1}^N viterbi[s, T] * a_{s,q_F}$; termination step

return the backtrace path by following backpointers to states back in time from $backpointer[q_F, T]$

Intuition

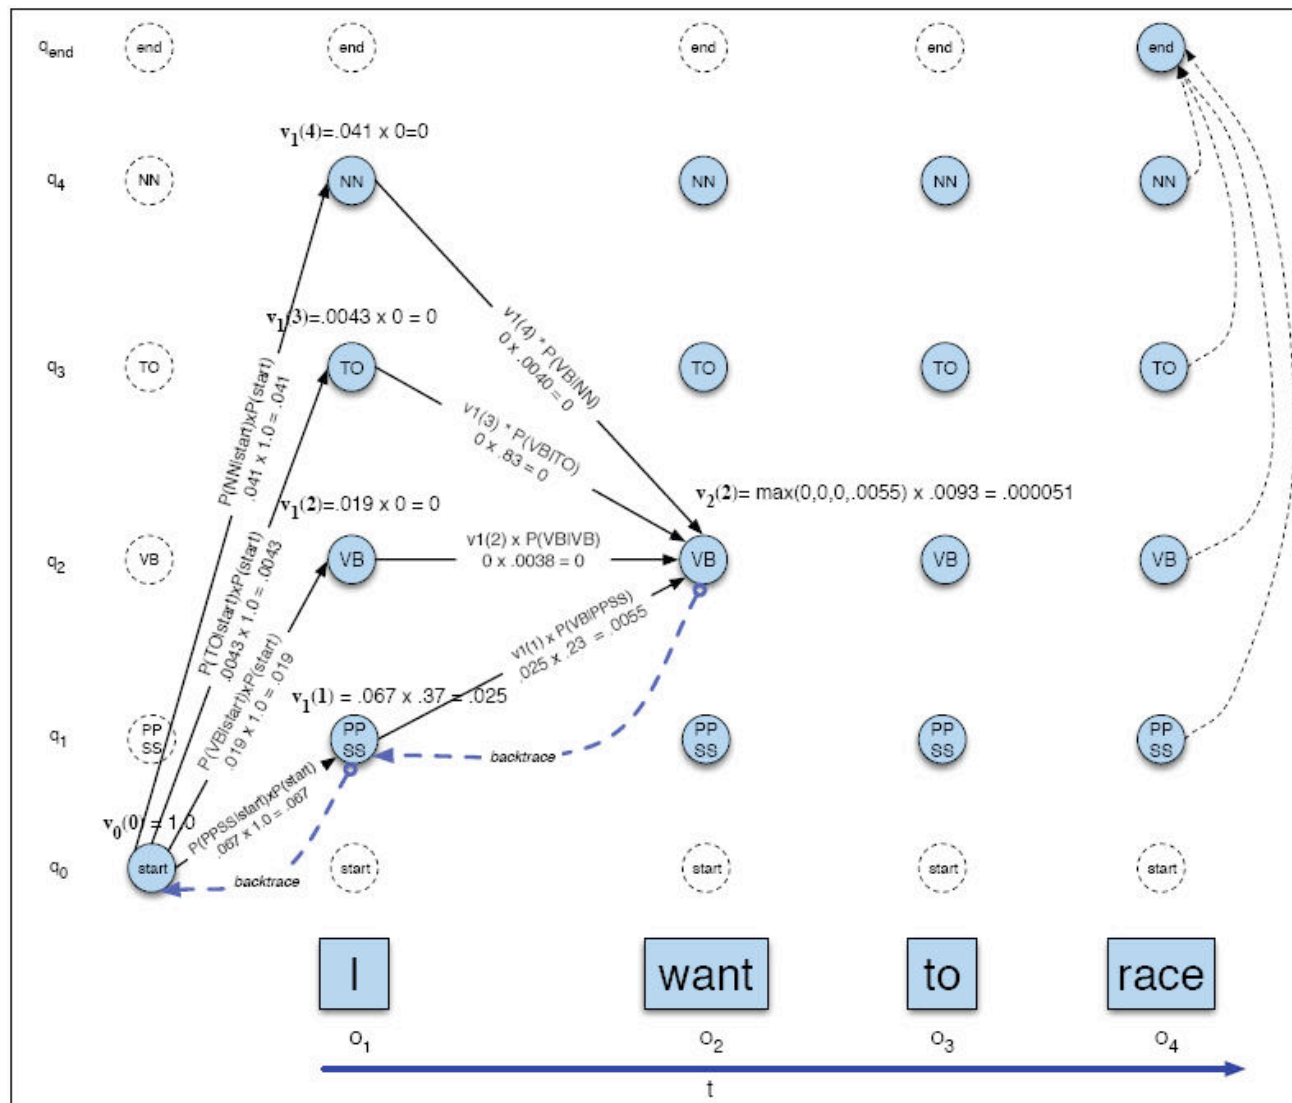
- The value in each cell is computed by taking the MAX over all paths that lead to this cell.

- $$v_t(j) = \max_{1 \leq i \leq N-1} v_{t-1}(i) a_{ij} b_j(o_t)$$

- An extension of a path from state i at time $t-1$ is computed by multiplying:

$v_{t-1}(i)$	the previous Viterbi path probability from the previous time step
a_{ij}	the transition probability from previous state q_i to current state q_j
$b_j(o_t)$	the state observation likelihood of the observation symbol o_t given the current state j

Viterbi example



Error Analysis: ESSENTIAL!!!

- Look at a confusion matrix

	IN	JJ	NN	NNP	RB	VBD	VBN
IN	—	.2			.7		
JJ	.2	—	3.3	2.1	1.7	.2	2.7
NN		8.7	—				.2
NNP	.2	3.3	4.1	—	.2		
RB	2.2	2.0	.5		—		
VBD		.3	.5			—	4.4
VBN		2.8				2.6	—

- See what errors are causing problems
 - ♦ Noun (NN) vs ProperNoun (NN) vs Adj (JJ)
 - ♦ Adverb (RB) vs Particle (RP) vs Prep (IN)
 - ♦ Preterite (VBD) vs Participle (VBN) vs Adjective (JJ)

Evaluation

- The result is compared with a manually coded “Gold Standard”
 - ◆ Typically accuracy reaches 96-97%
 - ◆ This may be compared with result for a baseline tagger (one that uses no context).
- Important: 100% is impossible even for human annotators.

Baseline

- Most frequent class baseline

Summary

- Part of speech tagging plays important role in TTS
- Most algorithms get 96-97% tag accuracy
- Not a lot of studies on whether remaining error tends to cause problems in TTS

Summary

I. Text Processing

1) Text Normalization

- Tokenization
- End of sentence detection
 - Methodology: decision trees

2) Homograph disambiguation

3) Part-of-speech tagging

- Methodology: Hidden Markov Models

II. Phonetic Analysis



Converting from words to phones

- Most important: dictionary

Dictionaries

- CMU dictionary: 127K words
 - ♦ <http://www.speech.cs.cmu.edu/cgi-bin/cmudict>

<i>ANTECEDENTS</i>	AE2 N T IH0 S IY1 D AH0 N T S	<i>PAKISTANI</i>	P AE2 K IH0 S T AE1 N IY0
<i>CHANG</i>	CH AE1 NG	<i>TABLE</i>	T EY1 B AH0 L
<i>DICTIONARY</i>	D IH1 K SH AH0 N EH2 R IY0	<i>TROTSKY</i>	T R AA1 T S K IY2
<i>DINNER</i>	D IH1 N ER0	<i>WALTER</i>	W AO1 L T ER0
<i>LUNCH</i>	L AH1 N CH	<i>WALTZING</i>	W AO1 L T S IH0 NG
<i>MCFARLAND</i>	M AH0 K F AA1 R L AH0 N D	<i>WALTZING(2)</i>	W AO1 L S IH0 NG

- Unisyn dictionary

going: { g * ou }.> i ng >

antecedents: { * a n . t ^ i . s ~ ii . d n ! t }> s >

dictionary: { d * i k . sh @ . n ~ e . r ii }

Lexicons and Lexical Entries

- In Festival you can explicitly give pronunciations for words

- ♦ Each lg/dialect has its own lexicon
- ♦ You can lookup words with
 - `(lex.lookup WORD)`
- ♦ You can add entries to the current lexicon
 - `(lex.add.entry NEWENTRY)`
- ♦ Entry: `(WORD POS (SYL0 SYL1...))`
- ♦ Syllable: `((PHONE0 PHONE1 ...) STRESS)`
- ♦ Example:

```
' ("cepstra" n ((k eh p) 1) ((s t r aa)
0) )))
```

Dictionaries aren't always sufficient

- Unknown words
 - ♦ Seem to be linear with number of words in unseen text
 - ♦ Mostly person, company, product names
 - ♦ But also foreign words, etc.
 - ♦ From a Black et al analysis
 - Of 39K tokens in part of the Wall Street Journal
 - 1775 (4.6%) were not in the OALD dictionary:

Names	Unknown	Typos and other
1360	351	64
76.6%	19.8%	3.6%

- So commercial systems have 3-part system:
 - ♦ Big dictionary
 - ♦ Special code for handling names
 - ♦ Machine learned LTS system for other unknown words

Names

- Big problem area is names
- Names are common
 - ♦ 20% of tokens in typical newswire text will be names
 - ♦ Spiegel (2003) estimate of US names:
 - 2 million surnames
 - 100,000 first names
 - ♦ Personal names: McArthur, D'Angelo, Jiminez, Rajan, Raghavan, Sondhi, Xu, Hsu, Zhang, Chang, Nguyen
 - ♦ Company/Brand names: Infinit, Kmart, Cytoc, Medamicus, Inforte, Aeon, Idexx Labs, Bebe

Names

- Methods:
 - ♦ Can do morphology (Walters -> Walter, Lucasville)
 - ♦ Can write stress-shifting rules (Jordan -> Jordanian)
 - ♦ Rhyme analogy: Plotsky by analogy with Trotsky (replace tr with pl)
 - ♦ Liberman and Church: for 250K most common names, got 212K (85%) from these modified-dictionary methods, used LTS for rest.
 - ♦ Can do automatic country detection (from letter trigrams) and then do country-specific rules

Letter-to-Sound Rules

- Festival LTS rules:
- (LEFTCONTEXT [ITEMS] RIGHTCONTEXT = NEWITEMS)
- Example:
 - ♦ (# [c h] C = k)
 - ♦ (# [c h] = ch)
- # denotes beginning of word
- C means all consonants
- Rules apply in order
 - ♦ “christmas” pronounced with [k]
 - ♦ But word with ch followed by non-consonant pronounced [ch]
 - E.g., “choice”

What about stress: practice

- Generally
- Pronounced
- Exception
- Dictionary
- Significant
- Prefix
- Exhale
- Exhalation
- Sally

Stress rules in LTS

- English famously evil: one from Allen et al 1987
- $V \rightarrow [1\text{-stress}] / X_C^* \{V_{\text{short}} C C? | V\} \{[V_{\text{short}} C^* | V]\}$
- Where X must contain all prefixes:
- Assign 1-stress to the vowel in a syllable preceding a weak syllable followed by a morpheme-final syllable containing a short vowel and 0 or more consonants (e.g. difficult)
- Assign 1-stress to the vowel in a syllable preceding a weak syllable followed by a morpheme-final vowel (e.g. oregano)
- etc

Modern method: Learning LTS rules automatically

- Induce LTS from a dictionary of the language
- Black et al. 1998
- Applied to English, German, French
- Two steps: alignment and (CART-based) rule-induction

Alignment

- Letters: c h e c k e d
- Phones: ch _ eh _ k _ t
- Black et al Method 1:

L:	c	a	k	e
P:	K	EY	K	ε

- ♦ First scatter epsilons in all possible ways to cause letters and phones to align
- ♦ Then collect stats for $P(\text{phone}|\text{letter})$ and select best to generate new stats

$$p(p_i|l_j) = \frac{\text{count}(p_i, l_j)}{\text{count}(l_j)}$$

- ♦ This iterated a number of times until settles (5-6)
- ♦ This is EM (expectation maximization) alg

Alignment

- Black et al method 2
- Hand specify which letters can be rendered as which phones
 - ♦ C goes to k/ch/s/sh
 - ♦ W goes to w/v/f, etc
 - ♦ An actual list:

c: k ch s sh t-s ε

e: ih iy er ax ah eh ey uw ay ow y-uw oy aa ε

- Once mapping table is created, find all valid alignments, find $p(\text{letter}|\text{phone})$, score all alignments, take best

Alignment

- Some alignments will turn out to be really bad.
- These are just the cases where pronunciation doesn't match letters:
 - ♦ Dept d ih p aa r t m ah n t
 - ♦ CMU s iy eh m y uw
 - ♦ Lieutenant l eh f t eh n ax n t (British)
- Also foreign words
- These can just be removed from alignment training


Building CART trees

- Build a CART tree for each letter in alphabet (26 plus accented) using context of +-3 letters
- # # # c h e c -> ch
- c h e c k e d -> _
- This produces 92-96% correct LETTER accuracy (58-75 word acc) for English

Improvements

- Take names out of the training data
- And acronyms
- Detect both of these separately
- And build special-purpose tools to do LTS for names and acronyms

Add more features



Summary

- Overview of Festival
 - ♦ Where it lives, its components
 - ♦ Its scripting language: Scheme
- Finishing up Part of Speech Tagging
- Phonetic Analysis
 - ♦ Dictionaries
 - ♦ Names
 - ♦ Letter-to-Sound Rules
 - (or “Grapheme-to-Phoneme Conversion”)